

Decision Trees on amazon reviews

June 30, 2018

```
In [1]: import sqlite3
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import plotly.plotly as py
import plotly.graph_objs as go
from sklearn.cross_validation import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from sklearn.cross_validation import cross_val_score
from collections import Counter
from sklearn.metrics import accuracy_score
from sklearn import cross_validation
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.metrics import confusion_matrix
from sklearn import metrics
from sklearn.metrics import roc_curve, auc
from nltk.stem.porter import PorterStemmer
from sklearn.decomposition import TruncatedSVD
from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import GridSearchCV
from sklearn.linear_model import LogisticRegression
```

```
/usr/local/lib/python3.6/dist-packages/sklearn/cross_validation.py:41: DeprecationWarning:
```

```
This module was deprecated in version 0.18 in favor of the model_selection module into which all
```

1 Loading the data

```
In [0]: final_data = pd.read_csv("final.csv")

final_data = final_data.drop(["Text"], axis = 1)
final_data = final_data.drop(final_data.columns[0], axis = 1)
```

2 Sorting and train/test split

```
In [0]: labels = final_data.Score
        final_data = final_data.sort_values("Time")
        final_data.shape

In [0]: n = final_data.shape[0]
        train_size = 0.7

        train_set = final_data.iloc[:int(n*train_size)]
        test_set = final_data.iloc[int(n*train_size):]

        X_train = train_set.CleanedText
        y_train = train_set.Score

        X_test = test_set.CleanedText
        y_test = test_set.Score
```

3 Cleaning the data

```
In [7]: import nltk
        nltk.download('stopwords')
        import re
        # Tutorial about Python regular expressions: https://pymotw.com/2/re/
        import string
        from nltk.corpus import stopwords
        from nltk.stem import PorterStemmer
        from nltk.stem.wordnet import WordNetLemmatizer

        stop = set(stopwords.words('english')) #set of stopwords
        sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer

        def cleanhtml(sentence): #function to clean the word of any html-tags
            cleanr = re.compile('<.*?>')
            cleantext = re.sub(cleanr, ' ', sentence)
            return cleantext
        def cleanpunc(sentence): #function to clean the word of any punctuation or special char
            cleaned = re.sub(r'[? ! | \\' " | #]', '', sentence)
            cleaned = re.sub(r'[, | | | ( | \\' / ]', ' ', cleaned)
            return cleaned
        print(stop)
        print('*****')
        print(sno.stem('tasty'))

[nltk_data] Downloading package stopwords to /content/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
{'am', 'below', 'through', 'mustn', 'mightn', 'itself', "didn't", 'm', 'the', 'haven', 'couldn
*****
```

tasti

4 Creating Word2vec model

```
In [0]: import gensim
        i=0
        train_sent=[]
        for sent in X_train:
            filtered_sentence=[]
            sent=cleanhtml(sent)
            for w in sent.split():
                for cleaned_words in cleanpunc(w).split():
                    if(cleaned_words.isalpha()):
                        filtered_sentence.append(cleaned_words.lower())
                    else:
                        continue
            train_sent.append(filtered_sentence)

In [0]: test_sent=[]
        for sent in X_test:
            filtered_sentence=[]
            sent=cleanhtml(sent)
            for w in sent.split():
                for cleaned_words in cleanpunc(w).split():
                    if(cleaned_words.isalpha()):
                        filtered_sentence.append(cleaned_words.lower())
                    else:
                        continue
            test_sent.append(filtered_sentence)

In [0]: from gensim.models import Word2Vec
        from gensim.models import KeyedVectors

        w2v_model = gensim.models.Word2Vec(train_sent,min_count=5,size=50, workers=4)
```

4.1 Average Word2Vec

```
In [12]: #AVG-W2V
        sent_vectors = []; # the avg-w2v for each sentence/review is stored in this list
        for sent in train_sent: # for each review/sentence
            sent_vec = np.zeros(50) # as word vectors are of zero length
            cnt_words = 0; # num of words with a valid vector in the sentence/review
```

```

    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            cnt_words = 1
        pass
    sent_vec /= cnt_words
    sent_vectors.append(sent_vec)
print(len(sent_vectors))
print(len(sent_vectors[0]))

```

254919

50

```

In [13]: sent_vectors2 = []; # the avg-w2v for each sentence/review is stored in this list
for sent in test_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    cnt_words = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            sent_vec += vec
            cnt_words += 1
        except:
            cnt_words = 1
        pass
    sent_vec /= cnt_words
    sent_vectors2.append(sent_vec)
print(len(sent_vectors2))
print(len(sent_vectors2[0]))

```

109252

50

```

In [0]: X_train1 = sent_vectors
        X_test1 = sent_vectors2

```

5 Applying DecisionTree Classifier

```

In [17]: from sklearn.tree import DecisionTreeClassifier

```

```

depth = list(range(2,10))

```

```

# empty list that will hold cv scores
cv_scores = []
my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train1)]
# perform 10-fold cross validation
for d in depth:
    dt = DecisionTreeClassifier(max_depth = d, min_samples_split = 1000, min_samples_
    scores = cross_val_score(dt, X_train1, y_train, cv = my_cv, scoring='accuracy')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

# determining best k
optimal_d = depth[MSE.index(min(MSE))]
print('\nThe optimal depth of the tree is %d.' % optimal_d)


# plot misclassification error vs k
plt.plot(depth, MSE)

for xy in zip(depth, np.round(MSE,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

plt.xlabel('Tree depth')
plt.ylabel('Misclassification Error')
plt.show()

```

The optimal depth of the tree is 8.



```
In [18]: dt = DecisionTreeClassifier(max_depth = 8, min_samples_split = 1000, min_samples_leaf
      dt.fit(X_train1,y_train)
      pred = dt.predict(X_test1)

      acc = accuracy_score(y_test, pred, normalize=True) * float(100)
      x = dt.predict(X_train1)
      tr_acc = accuracy_score(y_train, x, normalize=True) * float(100)
      print('\n****Train accuracy for k = {} is {:.2f}'.format(8,tr_acc))
      print('\n****Test accuracy for k = {} is {:.2f}'.format(8,acc))
```

```
****Train accuracy for k = 8 is 87.32
```

```
****Test accuracy for k = 8 is 85.31
```

6 Word2Vec-Tfidf

6.0.1 Note : Performed with sampled 50k datapoints

```
In [12]: from sklearn.feature_extraction.text import TfidfTransformer
      from sklearn.feature_extraction.text import TfidfVectorizer
```

```

tf_idf_vect = TfidfVectorizer()
final_tf_idf = tf_idf_vect.fit_transform(X_train)
tfidf_feat = tf_idf_vect.get_feature_names()
# tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

train_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in train_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model.wv[word]
            # obtain the tf_idf of a word in a sentence/review
            tfidf = final_tf_idf[row, tfidf_feat.index(word)]
            sent_vec += (vec * tfidf)
            weight_sum += tfidf

        except:
            weight_sum = 1
            pass

    sent_vec /= weight_sum
    #print(np.isnan(np.sum(sent_vec)))

    train_vectors.append(sent_vec)
    row += 1
print(len(train_vectors))
print(len(train_vectors[0]))

```

35000

50

```

In [13]: final_tf_idf = tf_idf_vect.fit_transform(X_test)
tfidf_feat = tf_idf_vect.get_feature_names()
# tfidf words/col-names
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val = tfidf

test_vectors = []; # the tfidf-w2v for each sentence/review is stored in this list
row=0;
for sent in test_sent: # for each review/sentence
    sent_vec = np.zeros(50) # as word vectors are of zero length
    weight_sum = 0; # num of words with a valid vector in the sentence/review

```

```

    for word in sent: # for each word in a review/sentence
        try:
            vec = w2v_model2.wv[word]
            # obtain the tf_idf of a word in a sentence/review
            tfidf = final_tf_idf[row, tfidf_feat.index(word)]
            sent_vec += (vec * tfidf)
            weight_sum += tfidf

        except:
            weight_sum = 1
            pass

    sent_vec /= weight_sum
    #print(np.isnan(np.sum(sent_vec)))

    test_vectors.append(sent_vec)
    row += 1
print(len(test_vectors))
print(len(test_vectors[0]))

```

15000

50

```

In [0]: X_train2 = train_vectors
        X_test2 = test_vectors

```

7 Applying DecisionTree Classifier

```

In [15]: from sklearn.tree import DecisionTreeClassifier

```

```

depth = list(range(2,15))

# empty list that will hold cv scores
cv_scores = []
my_cv = [(train,test) for train, test in TimeSeriesSplit(n_splits=10).split(X_train2)]
# perform 10-fold cross validation
for d in depth:
    dt = DecisionTreeClassifier(max_depth = d, min_samples_split = 500, min_samples_leaf = 10)
    scores = cross_val_score(dt, X_train2, y_train, cv = my_cv, scoring='accuracy')
    cv_scores.append(scores.mean())

# changing to misclassification error
MSE = [1 - x for x in cv_scores]

```



```

# determining best k
optimal_d = depth[MSE.index(min(MSE))]
print('\nThe optimal depth of the tree is %d.' % optimal_d)


# plot misclassification error vs k
plt.plot(depth, MSE)

for xy in zip(depth, np.round(MSE,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Tree depth')
plt.ylabel('Misclassification Error')
plt.show()

```

The optimal depth of the tree is 6.



```

In [17]: dt = DecisionTreeClassifier(max_depth = 6, min_samples_split = 500, min_samples_leaf = 500)
dt.fit(X_train2,y_train)
pred = dt.predict(X_test2)

acc = accuracy_score(y_test, pred, normalize=True) * float(100)
x = dt.predict(X_train2)
tr_acc = accuracy_score(y_train, x, normalize=True) * float(100)

```

```
print('\n****Train accuracy for k = {} is {:.2f}'.format(6,tr_acc))
print('\n****Test accuracy for k = {} is {:.2f}'.format(6,acc))
```

```
****Train accuracy for k = 6 is 86.09
```

```
****Test accuracy for k = 6 is 83.13
```

8 Conclusion

-

8.0.1 Average word2vec gives a training accuracy of 87.3 % and testing accuracy of 85.3 % with best depth of tree being '8'.

-

8.0.2 Tf-idf word2vec gives a training accuracy of 86 % and testing accuracy of 83.1 % with best depth of tree being '6'.